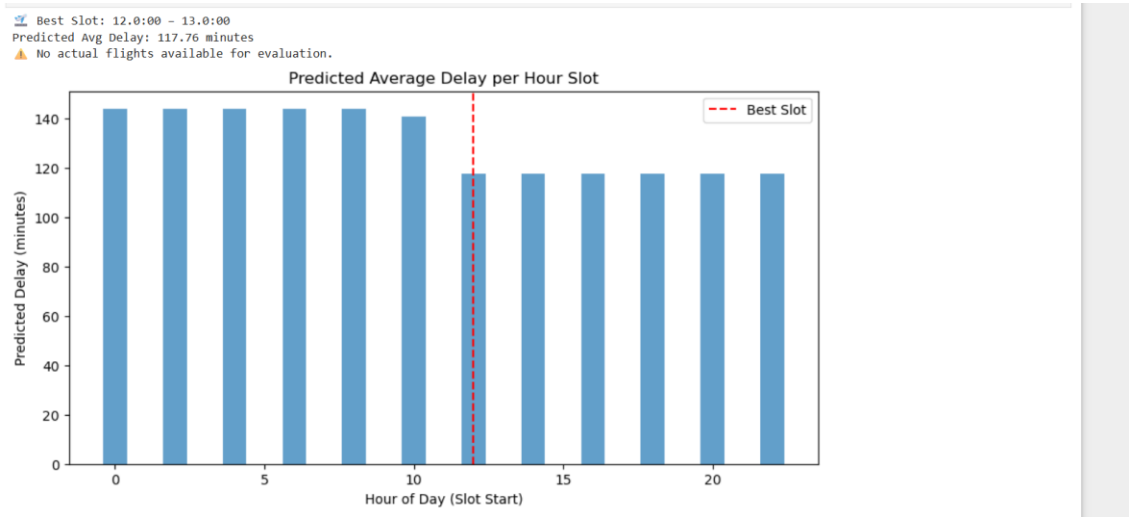


ML MODEL -RESULTS

METRICS:

	Model	MAE	RMSE	R2
0	RandomForest	20.201646	37.966809	0.581613
2	XGBoost	20.827460	37.075785	0.601020
1	GradientBoosting	21.050305	37.927764	0.582473
3	Ridge	31.976122	51.527597	0.229363
4	NeuralNet	32.015345	52.439434	0.201847

PREDICTIONS



EXPORT ML MODEL:

```
import os

os.makedirs("models", exist_ok=True) # create folder if not exists
joblib.dump(best_pipe, "models/best_pipe.pkl")

['models/best_pipe.pkl']
```

CODE:

```

import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error
from sklearn.preprocessing import LabelEncoder

# === 0. Load / parse your dataframe (replace with your file) ===
# df = pd.read_csv("flights_clean.csv") # or however you load it
# For safety, ensure columns exist and datetimes parse
df["STD"] = pd.to_datetime(df["STD"], errors="coerce")
df["ATD"] = pd.to_datetime(df["ATD"], errors="coerce")
df["Date"] = pd.to_datetime(df["Date"], errors="coerce")

# Drop rows with essential missing values
df = df.dropna(subset=["STD", "ATD", "Date", "Aircraft", "From", "To"])

# === 1. Feature engineering ===
df["ScheduledHour"] = df["STD"].dt.hour
df["Weekday"] = df["Date"].dt.weekday

# Create route
df["Route"] = df["From"].astype(str) + "-" + df["To"].astype(str)

# Target: departure delay in minutes (float)
df["DepartureDelay"] = (df["ATD"] - df["STD"]).dt.total_seconds() / 60.0

# If you want to remove extreme outliers, consider clipping:
# df = df[(df["DepartureDelay"] > -60) & (df["DepartureDelay"] < 24*60)]

# === 2. Encode categorical features ===

```

```

le_aircraft = LabelEncoder()
le_route = LabelEncoder()

df["AircraftEncoded"] = le_aircraft.fit_transform(df["Aircraft"])
df["RouteEncoded"] = le_route.fit_transform(df["Route"])

# Save mappings for safe transform on unseen labels
aircraft_map = {lab: idx for idx, lab in enumerate(le_aircraft.classes_)}
route_map = {lab: idx for idx, lab in enumerate(le_route.classes_)}

# === 3. Prepare X, y and train/test split ===
feature_cols = ["ScheduledHour", "Weekday", "AircraftEncoded", "RouteEncoded"]
X = df[feature_cols].copy()
y = df["DepartureDelay"].copy()

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# === 4. Train model ===
model = RandomForestRegressor(n_estimators=200, random_state=42, n_jobs=-1)
model.fit(X_train, y_train)

# Evaluate quickly
y_pred = model.predict(X_test)
print("MAE on test set: {:.2f} minutes".format(mean_absolute_error(y_test, y_pred)))

# === 5. Robust function to suggest best hour ===
def safe_encode_aircraft(a):
    # return encoded integer or -1 for unseen
    return aircraft_map.get(a, -1)

def safe_encode_route(from_code, to_code):

```

```

route = f"{from_code}-{to_code}"
return route_map.get(route, -1)

```

```

def suggest_best_time_for_flight(aircraft, origin, dest, weekday, candidate_hours=range(0,24)):

```

```

    """

```

```

    aircraft: string like "A320"

```

```

    origin, dest: strings e.g. "HYD", "DEL"

```

```

    weekday: int 0=Monday .. 6=Sunday (or use same convention as training)

```

```

    candidate_hours: iterable of hours to test (0..23)

```

```

    Returns (best_hour, predicted_delay_at_best_hour, df_with_all)

```

```

    """

```

```

    aircraft_enc = safe_encode_aircraft(aircraft)

```

```

    route_enc = safe_encode_route(origin, dest)

```

```

    if aircraft_enc == -1:

```

```

        print(f"Warning: aircraft '{aircraft}' not seen in training -> using -1 encoding.")

```

```

    if route_enc == -1:

```

```

        print(f"Warning: route '{origin}-{dest}' not seen in training -> using -1 encoding.")

```

```

    rows = []

```

```

    for h in candidate_hours:

```

```

        rows.append({

```

```

            "ScheduledHour": h,

```

```

            "Weekday": int(weekday),

```

```

            "AircraftEncoded": aircraft_enc,

```

```

            "RouteEncoded": route_enc

```

```

        })

```

```

    X_try = pd.DataFrame(rows, columns=feature_cols)

```

```

    preds = model.predict(X_try)

```

```

    X_try["PredictedDelay"] = preds

```

```

    best_idx = X_try["PredictedDelay"].idxmin()

```

```

    best_row = X_try.loc[best_idx]

```

```

return int(best_row["ScheduledHour"]), float(best_row["PredictedDelay"]), X_try

# === 6. Example usage ===
# Example: a flight HYD -> DEL on Wednesday (weekday=2), aircraft "A320"
best_hour, best_pred_delay, candidates_df = suggest_best_time_for_flight(
    aircraft="A320",
    origin="HYD",
    dest="DEL",
    weekday=2
)

print("Best hour:", best_hour, "Predicted delay (min):", round(best_pred_delay, 2))
# candidates_df holds predicted delays for all 24 hours if you want to inspect/plot

import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Models
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.linear_model import Ridge
from sklearn.neural_network import MLPRegressor
import xgboost as xgb

# === Data preparation ===
df["STD"] = pd.to_datetime(df["STD"], errors="coerce")
df["ATD"] = pd.to_datetime(df["ATD"], errors="coerce")

```

```

df["Date"] = pd.to_datetime(df["Date"], errors="coerce")

df = df.dropna(subset=["STD", "ATD", "Date", "Aircraft", "From", "To"])
df["ScheduledHour"] = df["STD"].dt.hour
df["Weekday"] = df["Date"].dt.weekday
df["Route"] = df["From"].astype(str) + "-" + df["To"].astype(str)
df["DepartureDelay"] = (df["ATD"] - df["STD"]).dt.total_seconds() / 60.0

# Features
cat_cols = ["Aircraft", "Route"]
num_cols = ["ScheduledHour", "Weekday"]
target = "DepartureDelay"

X = df[cat_cols + num_cols]
y = df[target]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# Preprocessing: OneHotEncoder for categorical vars
preprocessor = ColumnTransformer(
    transformers=[
        ("cat", OneHotEncoder(handle_unknown="ignore"), cat_cols),
        ("num", "passthrough", num_cols)
    ]
)

# === Define candidate models ===
models = {
    "RandomForest": RandomForestRegressor(n_estimators=200, random_state=42, n_jobs=-1),

```

```

"GradientBoosting": GradientBoostingRegressor(n_estimators=300, random_state=42),
"XGBoost": xgb.XGBRegressor(
    n_estimators=300, learning_rate=0.1, max_depth=6, random_state=42, n_jobs=-1
),
"Ridge": Ridge(alpha=1.0),
"NeuralNet": MLPRegressor(hidden_layer_sizes=(64,32), max_iter=500, random_state=42)
}

```

```

results = []

```

```

# === Train & Evaluate ===

```

```

for name, model in models.items():

```

```

    pipe = Pipeline(steps=[("pre", preprocessor), ("model", model)])

```

```

    pipe.fit(X_train, y_train)

```

```

    y_pred = pipe.predict(X_test)

```

```

    mae = mean_absolute_error(y_test, y_pred)

```

```

    rmse = np.sqrt(mean_squared_error(y_test, y_pred)) # ✅ works in all versions

```

```

    r2 = r2_score(y_test, y_pred)

```

```

    results.append({

```

```

        "Model": name,

```

```

        "MAE": round(mae, 2),

```

```

        "RMSE": round(rmse, 2),

```

```

        "R2": round(r2, 3)

```

```

    })

```

```

results_df = pd.DataFrame(results).sort_values(by="MAE")

```

```

print(results_df)

```

```

#select best model

```

```

best_model_name = results_df.iloc[0]["Model"]

```

```
best_model = models[best_model_name]

best_pipe = Pipeline(steps=[("pre", preprocessor), ("model", best_model)])


best_pipe.fit(X, y) # fit on full dataset
```

```
import pandas as pd

import numpy as np

from sklearn.metrics import mean_absolute_error, mean_squared_error

import matplotlib.pyplot as plt
```

```
#  Suggest best slot across a range
```

```
def suggest_best_time_slots(aircraft, route, weekday, slot_size, best_pipe, df):
```

```
    times = range(0, 24) # scan full day
```

```
    predictions = []
```

```
    # 1. Predict delay for each hour
```

```
    for t in times:
```

```
        test_input = pd.DataFrame([
```

```
            "Aircraft": aircraft,
```

```
            "Route": route,
```

```
            "Weekday": weekday,
```

```
            "ScheduledHour": t
```

```
        ]])
```

```
        pred = best_pipe.predict(test_input)[0]
```

```
        predictions.append((t, pred))
```

```
pred_df = pd.DataFrame(predictions, columns=["Hour", "PredictedDelay"])
```

```
# 2. Group into slots (e.g. 6–7, 7–8 ...)
```

```
pred_df["SlotStart"] = (pred_df["Hour"] // slot_size) * slot_size
```

```
slot_summary = pred_df.groupby("SlotStart")["PredictedDelay"].mean().reset_index()
```



```
# 3. Find best slot (minimum avg delay)
```

```
best_slot = slot_summary.loc[slot_summary["PredictedDelay"].idxmin()]
```

```
# 4. Evaluate efficiency on actual data
```

```
mask = (
```

```
    (df["Aircraft"] == aircraft) &
```

```
    (df["Route"] == route) &
```

```
    (df["Weekday"] == weekday)
```

```
)
```

```
actual_data = df.loc[mask, ["ScheduledHour", "DepartureDelay", "Aircraft", "Route", "Weekday"]]
```

```
if not actual_data.empty:
```

```
    pred_delays = best_pipe.predict(actual_data[["Aircraft", "Route", "Weekday", "ScheduledHour"]])
```

```
    actual_data = actual_data.assign(PredictedDelay=pred_delays)
```

```
    mae = mean_absolute_error(actual_data["DepartureDelay"], actual_data["PredictedDelay"])
```

```
    rmse = np.sqrt(mean_squared_error(actual_data["DepartureDelay"],  
actual_data["PredictedDelay"]))
```

```
else:
```

```
    mae, rmse = None, None
```

```
return best_slot, slot_summary, actual_data, mae, rmse
```

```
#  Example usage
```

```
best_slot, slot_summary, slot_df, mae, rmse = suggest_best_time_slots(
```

```
    df=df,
```

```
    aircraft="A320",
```

```
    route="MUM-DEL",
```

```
    weekday=0,    # Monday
```

```
    slot_size=2,  # 1-hour slots
```

```
    best_pipe=best_pipe
```

)

```
print(f"📅 Best Slot: {best_slot['SlotStart']}:00 – {best_slot['SlotStart']+1}:00")
```

```
print(f"Predicted Avg Delay: {best_slot['PredictedDelay']:.2f} minutes")
```

if mae is not None:

```
    print(f"📊 Model Efficiency → MAE={mae:.2f}, RMSE={rmse:.2f}")
```

else:

```
    print("⚠️ No actual flights available for evaluation.")
```

✅ Plot all slots

```
plt.figure(figsize=(10,5))
```

```
plt.bar(slot_summary["SlotStart"], slot_summary["PredictedDelay"], width=0.8, alpha=0.7)
```

```
plt.axvline(best_slot["SlotStart"], color="red", linestyle="--", label="Best Slot")
```

```
plt.xlabel("Hour of Day (Slot Start)")
```

```
plt.ylabel("Predicted Delay (minutes)")
```

```
plt.title("Predicted Average Delay per Hour Slot")
```

```
plt.legend()
```

```
plt.show()
```